

Monitoring Neo4J

Monitoring in Neo4J

- Authentication events are written to the security.log file
- Event categories that we monitor in log files include:
 - Queries
 - Transactions
 - Connections
 - Memory

Collecting metrics

- Can configure to collect metrics that are related to events
- Can be viewed in tools - such as Grafana
- Can configure a tool such as Nagios to provide alerts when certain metrics are detected in Neo4j




Monitoring queries

- Write information about queries to the query.log file
- Log information about queries that took a long time to complete
- Can also monitor currently running queries and if need be, kill them if they are taking too long.



Configuring query logging

- Should enable logging for queries
- Set a threshold for the length of time queries take
- Regularly inspect the query.log file
- Determine if queries are taking longer duration



Configuring query logging

- Examples:

- Options: [OFF, INFO, VERBOSE]
- `dbms.logs.query.enabled=VERBOSE`
- #If the execution of query takes more time than this threshold, the query is logged once completed
- `dbms.logs.query.threshold=1000ms`
- #Log parameters for the executed queries being logged
- `dbms.logs.query.parameter_logging_enabled=true`
- #Log detailed time information for the executed queries being logged
- `dbms.logs.query.time_logging_enabled=true`

Configuring query logging

• Examples:

- #Log amount of total allocated bytes for the executed queries being logged. The logged number is cumulative over the duration of the query
- `dbms.logs.query.allocation_logging_enabled=true`
- #Log page hits and page faults for the executed queries being logged.
- `dbms.logs.query.page_logging_enabled=true`
- #Enables or disables tracking of how much time a query spends actively executing on the CPU
- `dbms.track_query_cpu_time=true`
- #Enables or disables tracking of how many bytes are allocated by the execution of a query.
- If enabled, calling `dbms.listQueries` will display the allocated bytes
- `dbms.track_query_allocation=true`

Dynamic settings

- All the above listed settings are dynamic.
- To view the list of settings which can be set dynamically:
 - `CALL dbms.listConfig()`
 - `YIELD name, dynamic`
 - `WHERE dynamic`
 - `RETURN name`
 - `ORDER BY name`
- To configure dynamic setting:
 - `CALL dbms.setConfigValue('dbms.logs.query.enabled', 'info')`

Read the logs

- `tail -f -n 50 /logs/query.log`

Create some nodes

- `cypher-shell -u neo4j -p secret -d neo4j --format plain`
 - `:use neo4j;`
 - `MATCH (n) DETACH DELETE n;`
 - `#Create some records`
 - `FOREACH (i IN RANGE(1,200) | CREATE (:Person {name:'Person' + i}));`

Exercise: Monitoring queries

- `cypher-shell -u neo4j -p secret -d neo4j --format plain`
 - #Execute a query that runs for longer than 1000 ms:
 - `MATCH (a), (b), (c), (d) RETURN count(id(a));`

Viewing currently running queries

\$



To enjoy the full Neo4j Browser experience, we advise you to use [Neo4j Browser Sync](#)

X

\$:queries



Database URI	User	Query	Params	Meta	Elapsed time	Kill
bolt://ec2-3-82-2-121.compu...	admin	CALL dbms.listQueries	{}	{}	0 ms	⊖
bolt://ec2-3-82-2-121.compu...	reader	MATCH (a), (b), (c), (d), (e) RETU...	{}	{ "type": "user-direct", "app": "n...	1344899 ms	⊖

Found 2 queries running on one server

AUTO-REFRESH ☐ OFF



Killing a long-running query

\$

To enjoy the full Neo4j Browser experience, we advise you to use [Neo4j Browser Sync](#)

\$:queries

Database URI	User	Query	Params	Meta	Elapsed time	Kill
bolt://ec2-3-82-2-121.compu...	admin	CALL dbms.listQueries	{}	{}	0 ms	⊖
bolt://ec2-3-82-2-121.compu...	reader	MATCH (a), (b), (c), (d), (e) RETU...	{}	{ "type": "user-direct", "app": "n...	1344899 ms	⊖ ✕

Found 2 queries running on one server

AUTO-REFRESH ☐ OFF

Exercise #1: Monitoring queries

- Modify the neo4j.conf file to create a log record if a query exceeds 1000 ms
- `vim /var/lib/neo4j/conf/neo4j.conf`
 - `dbms.logs.query.enabled=VERBOSE`
 - `dbms.logs.query.threshold=1000ms`
 - `dbms.logs.query.parameter_logging_enabled=true`
 - `dbms.logs.query.time_logging_enabled=true`
 - `dbms.logs.query.allocation_logging_enabled=true`
 - `dbms.logs.query.page_logging_enabled=true`
 - `dbms.track_query_cpu_time=true`
 - `dbms.track_query_allocation=true`
- `neo4j restart`

Exercise: Monitoring queries

- `cypher-shell -u neo4j -p secret -d neo4j --format plain`
 - #Execute a query that runs for longer than 1000 ms:
 - `MATCH (a), (b), (c), (d) RETURN count(id(a));`

Exercise: Monitoring queries

- Wait about a minute, it should complete.

```

Terminal
[ubuntu@ip-172-31-24-255:~$ /usr/bin/cypher-shell --format plain
[username: reader
[password: ****
[neo4j> MATCH (a), (b), (c), (d) RETURN count(id(a));
count(id(a))
855036081
neo4j> |

```

- View the query.log. Is there a record for this query?

```

[ubuntu@ip-172-31-24-255:/var/log/neo4j$ sudo systemctl restart neo4j
ubuntu@ip-172-31-24-255:/var/log/neo4j$ tail query.log
2019-02-15 18:33:48.069+0000 INFO 1929 ms: (planning: 1898, cpu: 1812, waiting: 0) - 260597464 B - 0 page hits, 0 page faults - embedded-session
- MATCH (a: ` This query is just used to load the cypher compiler during warmup. Please ignore `) RETURN a LIMIT 0 - {} - {}
2019-02-15 18:35:45.642+0000 INFO 26738 ms: (planning: 176, cpu: 26729, waiting: 0) - 26865112 B - 29415 page hits, 0 page faults - bolt-session bolt reader neo4j-java/dev client/127.0.0.1:39750 server/127.0.0.1:7687> reader - MATCH (a), (b), (c), (d) RETURN count(id(a)); - {} - {}
ubuntu@ip-172-31-24-255:/var/log/neo4j$ |

```


Exercise #1: Monitoring queries

- In cypher-shell session, enter query that will execute for an even longer time
 - MATCH (a), (b), (c), (d), (e) RETURN count(id(a));
- Open a new terminal window and log in to cypher-shell
- Execute the Cypher statement to list transactions
- Do you see the query?

**CALL dbms.listTransactions() yield username,
currentQueryId, currentQuery, elapsedTimeMillis;**

```

Terminal
ubuntu@ip-172-31-24-255:~$ /usr/bin/cypher-shell --format plain
username: admin
password: *****
neo4j> CALL dbms.listTransactions() yield username, currentQueryId, currentQuery, elapsedTimeMillis;
username, currentQueryId, currentQuery, elapsedTimeMillis
"reader", "query-4", "
MATCH (a), (b), (c), (d), (e) RETURN count(id(a));", 367299
"admin", "query-14", "CALL dbms.listTransactions() yield username, currentQueryId, currentQuery, elapsedTimeMillis;", 26
neo4j> |

```

Exercise #1: Monitoring queries

- Execute the Cypher statement to kill the long-running query.
 - `CALL dbms.listTransactions() yield username, currentQueryId, currentQuery, elapsedTimeMillis;`
 - `CALL dbms.killQuery('query-id');`

```
neo4j> CALL dbms.listTransactions() yield username, currentQueryId, currentQuery, elapsedTimeMillis;
username, currentQueryId, currentQuery, elapsedTimeMillis
"reader", "query-4", "
MATCH (a), (b), (c), (d), (e) RETURN count(id(a));", 367299
"admin", "query-14", "CALL dbms.listTransactions() yield username, currentQueryId, currentQuery, elap
[neo4j> CALL dbms.killQuery('query-4');
queryId, username, message
"query-4", "reader", "Query found"
neo4j> |
```

- Observe in other session that the query has been killed.

Automating monitoring of queries

- Automate the killing of long-running queries:
 - CALL dbms.listQueries() YIELD query, elapsedTimeMillis, queryId, username
 - WHERE NOT query CONTAINS toLower('LOAD')
 - AND elapsedTimeMillis > 1000
 - WITH query, collect(queryId) AS q
 - CALL dbms.killQueries(q) YIELD queryId
 - RETURN query, queryId;

Configuring transaction guard

- Example of lock acquisition timeout and transaction guard
- `vim /var/lib/neo4j/conf/neo4j.conf`
 - # transaction guard: max duration of any transaction
 - `dbms.transaction.timeout=1s`
 - # max time to acquire write lock
 - `dbms.lock.acquisition.timeout=10ms`
- `neo4j restart`
- Enter this statement to create multiple Person nodes:
 - `FOREACH (i IN RANGE(1,1000000) | CREATE (:Person {name:'Person' + i}));`
- Do you receive an error?
 - A record is written to the debug.log file

Configuring transaction guard

- View the record written to debug.log

```
ubuntu@ip-172-31-24-255:/var/log/neo4j$ tail -n 2 debug.log
2019-02-15 23:08:31.803+0000 INFO [o.n.i.d.DiagnosticsManager] --- SERVER STARTED END ---
2019-02-15 23:08:58.522+0000 WARN [o.n.k.i.a.t.m.KernelTransactionMonitor] Transaction KernelTransactionImplementationHandle{txReuseCou
nt=0, tx=KernelTransaction[0]} timeout.
ubuntu@ip-172-31-24-255:/var/log/neo4j$ |
```

Monitoring connections

- Can view the current connections to a Neo4j instance from cypher-shell

- Call `dbms.listConnections()`;

```
neo4j> call dbms.listConnections();
```

connectionId	connectTime	connector	username	userAgent	serverAddress	clientAddress
"bolt-97"	"2019-02-19T20:35:30.513Z"	"bolt"	"admin"	"neo4j-javascript/1.7.2"	"172.31.24.255:7687"	"216.45.71.93:57415"
"bolt-840"	"2019-02-19T21:58:04.65Z"	"bolt"	"publisher"	"neo4j-java/1.7.2-7165cef0d4b602da30b65613977744ad661c2e67"	"127.0.0.1:7687"	"127.0.0.1:40344"
"bolt-779"	"2019-02-19T21:52:05.978Z"	"bolt"	"admin"	"neo4j-java/dev"	"127.0.0.1:7687"	"127.0.0.1:40338"
"bolt-839"	"2019-02-19T21:58:04.358Z"	"bolt"	"publisher"	"neo4j-java/1.7.2-7165cef0d4b602da30b65613977744ad661c2e67"	"127.0.0.1:7687"	"127.0.0.1:40340"

```
4 rows available after 1 ms, consumed after another 0 ms
neo4j> |
```

- Terminate the connection

- `dbms.killConnection()`

```
neo4j> CALL dbms.killConnections(['bolt-1362','bolt-1363']);
```

connectionId	username	message
"bolt-1362"	"publisher"	"Connection found"
"bolt-1363"	"publisher"	"Connection found"

```
2 rows available after 14 ms, consumed after another 0 ms
neo4j> |
```

Logging HTTP requests

- You can set this property in neo4j.conf to log HTTP requests:
 - `vim /var/lib/neo4j/conf/neo4j.conf`
 - `dbms.logs.http.enabled=true`
 - `neo4j restart`
- View the records in the http.log file
 - `tail -f -n 10 /logs/http.log`

```
debug.log http.log query.log security.log
ubuntu@ip-172-31-24-255:/var/log/neo4j$ cat http.log
2019-02-20 14:41:07.614+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:07.609+0000] 216.45.71.93 - [Wed Feb 20 14:41:07 UTC 2019] "/?null" 303 -1 "" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 40
2019-02-20 14:41:07.697+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:07.697+0000] 216.45.71.93 - [Wed Feb 20 14:41:07 UTC 2019] "/browser/?null" 200 2895 "" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 10
2019-02-20 14:41:08.773+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:08.773+0000] 216.45.71.93 - [Wed Feb 20 14:41:08 UTC 2019] "/browser/main.chunkhash.bundle.js?null" 200 1650952 "http://ec2-54-158-36-157.compute-1.amazonaws.com:7474/browser/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 946
2019-02-20 14:41:11.042+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:11.042+0000] 216.45.71.93 - [Wed Feb 20 14:41:11 UTC 2019] "/browser/vendors-main.chunkhash.bundle.js?null" 200 3301516 "http://ec2-54-158-36-157.compute-1.amazonaws.com:7474/browser/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 3217
2019-02-20 14:41:12.503+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:12.502+0000] 216.45.71.93 - [Wed Feb 20 14:41:12 UTC 2019] "/?null" 200 232 "http://ec2-54-158-36-157.compute-1.amazonaws.com:7474/browser/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 57
2019-02-20 14:41:14.020+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:14.020+0000] 216.45.71.93 - [Wed Feb 20 14:41:14 UTC 2019] "/browser/assets/fonts/OpenSans-Semibold.ttf?null" 200 221328 "http://ec2-54-158-36-157.compute-1.amazonaws.com:7474/browser/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 3
2019-02-20 14:41:14.024+0000 INFO [REQUEST] [AsyncLog @ 2019-02-20 14:41:14.023+0000] 216.45.71.93 - [Wed Feb 20 14:41:14 UTC 2019] "/browser/assets/fonts/OpenSans-Light.ttf?null" 200 222412 "http://ec2-54-158-36-157.compute-1.amazonaws.com:7474/browser/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/72.0.3626.109 Safari/537.36" 2
ubuntu@ip-172-31-24-255:/var/log/neo4j$
```

Monitoring memory usage

- Default configurations in neo4j.conf are useful for small database
- In production, make sure that settings for the JVM are the best ones

Initial memory settings for database

- Obtain recommendation for settings related to memory:
 - `neo4j-admin memrec`
- Provides recommended memory settings

```
ubuntu@ip-172-31-24-255:/var/log/neo4j$ /usr/bin/neo4j-admin memrec --database=movie3.db
# Memory settings recommendation from neo4j-admin memrec:
#
# Assuming the system is dedicated to running Neo4j and has 7700m of memory,
# we recommend a heap size of around 3500m, and a page cache of around 1800m,
# and that about 2400m is left for the operating system, and the native memory
# needed by Lucene and Netty.
#
# Tip: If the indexing storage use is high, e.g. there are many indexes or most
# data indexed, then it might advantageous to leave more memory for the
# operating system.
#
# Tip: The more concurrent transactions your workload has and the more updates
# they do, the more heap memory you will need. However, don't allocate more
# than 31g of heap, since this will disable pointer compression, also known as
# "compressed oops", in the JVM and make less effective use of the heap.
#
# Tip: Setting the initial and the max heap size to the same value means the
# JVM will never need to change the heap size. Changing the heap size otherwise
# involves a full GC, which is desirable to avoid.
#
# Based on the above, the following memory settings are recommended:
dbms.memory.heap.initial_size=3500m
dbms.memory.heap.max_size=3500m
dbms.memory.pagecache.size=1800m
#
# The numbers below have been derived based on your current data volume in database and ind
# They can be used as an input into more detailed memory analysis.
# Lucene indexes: 0.0
# Data volume and native indexes: 54800k
```

Recommendation for log files

- Each type of log file should
 - Have its maximum size defined and
 - # of log files to keep
- # Number of HTTP logs to keep.
- `dbms.logs.http.rotation.keep_number=5`
- # Size of each HTTP log that is kept. (k,m,g)
- `dbms.logs.http.rotation.size=20m`
- # Number of query logs to keep.
- `dbms.logs.query.rotation.keep_number=5`
- # Size of each query log that is kept.
- `dbms.logs.query.rotation.size=20m`

Collecting metrics

- Automatically collects metrics
- Can disable them by setting
 - `metrics.enabled=false`
- Metrics are collected in CSV format
 - `cd /var/lib/neo4j/metrics`
 - `ls -al`